

Thanh-Long V. Le

[Homepage](#) [Google Scholar](#) [LinkedIn](#) [Email](#)

Research Interests

My research centers on **large language models** (LLMs), with a focus on enhancing their **reasoning** capabilities through **reinforcement learning** and other **post-training** techniques. Recently, I have also begun exploring **generative modeling**, particularly **diffusion models for video generation** and **diffusion-based LLMs**.

Education

Korea Advanced Institute of Science and Technology (KAIST AI) <i>M.S. in Artificial Intelligence</i>	Feb. 2025 – Feb. 2027 Seoul, South Korea
• GPA: 4.0 • Graduate Researcher at MLILAB	
Korea Advanced Institute of Science and Technology (KAIST) <i>B.S. in Computer Science and Artificial Intelligence</i>	Sep. 2019 – Feb. 2024 Daejeon, South Korea
• GPA: 3.98 • Undergraduate Researcher at NMSL • Graduated Magna Cum Laude (ranking in top 8% of the class) • Minor in Business and Technology Management	

Publications

No Prompt Left Behind: Exploiting Zero-Variance Prompts in LLM Reinforcement Learning via Entropy-Guided Advantage Shaping

Thanh-Long V. Le, Myeongho Jeon, Kim Vu, Viet Lai, Eunho Yang

Under review at *International Conference on Learning Representations (ICLR) 2026*

Design Opportunities for Explainable AI Paraphrasing Tools

Yewon Kim, Thanh-Long V. Le, Donghwi Kim, Mina Le, Sung-Ju Lee

Conference on Designing Interactive Systems (DIS) 2025

(FL)²: Overcoming Few Labels in Federated Semi-Supervised Learning

Seungjoo Lee, Thanh-Long V. Le, Jaemin Shin, Sung-Ju Lee

Conference on Neural Information Processing Systems (NeurIPS) 2024

Research Experience

Adobe Research

Research Collaborator, Advised by Dr. Viet Lai

Aug. 2025 – Present

Remote

- Led the RL-ZVP project, developing a novel algorithm that enhances LLM reinforcement learning by leveraging learning signals from zero-variance prompts.
- Currently working on improving existing RLVR techniques on large language diffusion models.

Machine Learning and Intelligence Lab, KAIST AI

Feb. 2025 – Present

Graduate Researcher, Advised by Prof. Eunho Yang

Seoul, South Korea

- Conducting research on LLM post-training, including LLM reasoning and alignment through reinforcement learning/preference optimization.
- Currently focusing on augmenting RLVR with various dense learning signal to improve the reasoning ability of LLM.
- Leading a collaborative project with AITRICS on developing a large reasoning model for highly imbalanced clinical event prediction using patient EHR data.

Networking and Mobile Systems Lab, KAIST

May. 2023 – Sep. 2024

Undergraduate Researcher, Advised by Prof. Sung-Ju Lee

Daejeon, South Korea

- Conducted research on natural language processing, federated learning and human-computer interaction.
- Actively participated in the Exphrase project – an AI-based paraphrasing system equipped with five supplementary features: AI Translation, AI Confidence Score, AI Explanation, Example Usages, Statistics.
- Addressed label deficiency in federated learning by implementing semi-supervised learning with adaptive thresholding and sharpness-aware regularization.

Work Experience

Luxoft – BMW Korea	Mar. 2024 – Dec. 2024
<i>C++ Software Engineer</i>	<i>Seoul, South Korea</i>
<ul style="list-style-type: none">Developed and optimized features for map provision application to enhance the accuracy and reliability of ADAS (Advanced Driving Assistance System) functionalities of BMW vehicles.Detected, analyzed, and prepared reports and fixes for vehicle anomalies by leveraging software logs and simulations.	
Coc Coc Company Ltd.	Sep. 2021 – Oct. 2022
<i>Software Engineer</i>	<i>Hanoi, Vietnam</i>
<ul style="list-style-type: none">Maintained the functionality of three key components of Coc Coc Search's backend system – compositor, knowledge graph, and search engine.Enhanced the quality of Coc Coc's Answer Box by implementing multilayer rule-based question-answering system and incorporating various alternative logic to address common query patterns, resulting in a 200% increase in impression and a 30% increase in click-through rate.Detected and fixed a critical issue in Coc Coc Translate which related to the time-to-live setting of cached translation results, resulting in a 5-time cost reduction for external translation API calls.Integrated Prometheus into services to enable automatic recording of metrics, eliminating the need for manual retrieval via logs.Conducted a full codebase refactoring of compositor system to eliminate potential bugs and technical debt.Utilized SQL to collect and analyze usage data of Coc Coc Translate over the course of a year, providing insights to assist managers in decision making.	

Projects

Improve the Robustness of QA System to Entity Renaming	Sep. 2022 – Dec. 2022
<ul style="list-style-type: none">Conducted experiments to evaluate and compare the robustness of BERT, RoBERTa, and SpanBERT against entity renaming on machine reading comprehension task across five distinct datasets.Proposed a novel improvement method named Anonymized Training with optional Anonymized InferenceAchieved a substantial mitigation in performance drop from over 30% to less than 5% for SpanBERT against entity renaming, with BERT and RoBERTa's performance drop being further reduced.	
Semi-Supervised Semantic Segmentation with Cross-Consistency Training	Sep. 2022 – Dec. 2022
<ul style="list-style-type: none">Replicate the findings presented in the original paper by Ouali et al, which focus on semi-supervised semantic segmentation using a shared encoder and multiple decoder models on the PASCAL VOC dataset.Conducted additional experiments with different encoder backbone, generative adversarial networks, and temporal ensembling, resulting in an improved mIoU score of 73.5 compared to the paper's result of 73.2.	

Contrastive Learning for Price Match Guarantee	Mar. 2022 – Jun. 2022
<ul style="list-style-type: none">Applied deep learning techniques to solve the product matching problem by utilizing image and title data.Implemented and integrated the triplet loss function into the training of ResNet18, resulting in a significant improvement of the model's validation F1 score from 60% to 71%.Achieved the highest validation F1 score of 72.46% for images matching and 68.27% for titles matching.	

Honors and Awards

Qualcomm-KAIST Innovation Award (2023): Winner.
KAIST Dean's List (2020): Top 3% students of the department.
KAIST International Undergraduate Scholarship (2019): Full-ride scholar.
Deakin Vice-Chancellor's International Scholarship (2019): Recipient.
Vietnam National Mathematical Olympiad (2018): Silver medal.

Technical Skills

Languages: C++, C, Python, Java, Go, SQL, HTML, CSS, JavaScript.
Technologies/Frameworks: PyTorch, verl, TRL, Unsloth, Jax, Haliax, Levanter, Linux, Git, Vim, Bazel, Bash Scripting.
Certificates: 8.0 in IELTS, 1500 in SAT, 800 in SAT Math Level 2, 800 in SAT Physics.